

Model Training Pipeline Documentation

Overview

This document provides a detailed description of the complete training pipeline for the Smart List Project - Product Recommendation Model. The model predicts the next product category a client is likely to purchase based on their historical purchase patterns and demographic information.

Table of Contents

1. Data Source and Identification
2. Data Preprocessing
3. Feature Engineering
4. Sequence Building
5. Train/Validation/Test Split
6. Model Training
7. Model Evaluation
8. Model Registration
9. Key Parameters

1. Data Source and Identification

1.1 Source Table

- **Table Name:** dl_tenants_daas.us_wealth_management.wealth_management_client_metrics
- **Platform:** Databricks Spark SQL
- **Data Type:** Client demographic data

1.2 Record Identification

The pipeline loads all records from the source table without initial filtering:

```
df_raw = spark.table("dl_tenants_daas.us_wealth_management.wealth_management_client_metrics")
```

2. Data Preprocessing

2.1 Product Category Transformation

The raw data contains multiple product classification fields:

- prod_lob (Product Line of Business)
- sub_product_level_1
- sub_product_level_2

These are consolidated into a single `product_category` field with the following categories:

- **LIFE_INSURANCE:** Life insurance products (VLI, WL, UL/IUL, TERM, etc.)
- **RETIREMENT:** Retirement products (401K, IRA, 403B, SEP, etc.)
- **INVESTMENT:** Investment products (Brokerage, Advisory, Direct)
- **NETWORK_PRODUCTS:** Network-related products
- **DISABILITY:** Disability insurance products
- **HEALTH:** Health insurance products
- **OTHER:** All other products

Transformation Logic: Hierarchical rules applied in order:

1. Check `prod_lob` first
2. Then check `sub_product_level_1`
3. Finally check `sub_product_level_2` with pattern matching

2.2 Data Filtering

2.2.1 Required Fields Filter

Keep only records with:

- Non-null `cont_id` (Client ID)

- Non-null `register_date` (Event date)
- Non-null `product_category` (After transformation)

Selected Columns:

- Client identifiers: `cont_id`
- Product information: `product_category`, `register_date`
- Financial metrics: `acct_val_amt`, `face_amt`, `cash_val_amt`, `wc_total_assets`
- Asset mix: `wc_assetmix_stocks`, `wc_assetmix_bonds`, `wc_assetmix_mutual_funds`, `wc_assetmix_annuity`, `wc_assetmix_deposits`, `wc_assetmix_other_assets`
- Demographics: `psn_age`, `client_seg`, `client_seg_1`, `aum_band`
- Channel information: `channel`, `agent_segment`, `branchoffice_code`
- Policy status: `policy_status`

2.2.2 Policy Status Filter

Filter to keep only **Active** policies:

```
df_events = df_events.filter(F.col("policy_status") == "Active")
```

2.2.3 Optional Sampling

For computational efficiency, apply random sampling:

- **Sample Fraction:** 0.2 (20% of data)
- **Seed:** 42 (for reproducibility)
- **Sampling Method:** Without replacement

Result: ~48,978,063 event rows after sampling

2.3 Event Ordering

Events are ordered chronologically per client:

1. Convert `register_date` to timestamp: `register_ts`
2. Create window function partitioned by `cont_id`, ordered by `register_ts`
3. Assign `event_idx` (row number) to each event per client

Purpose: Establish temporal sequence for each client's purchase history.

3. Feature Engineering

3.1 Product Vocabulary Building

Build a mapping between product categories and numeric IDs:

1. Extract all unique `product_category` values from the entire dataset
2. Sort alphabetically
3. Create mappings:
 - `prod2id`: Product name → ID (starting from 1, 0 reserved for padding)
 - `id2prod`: ID → Product name
4. **Vocabulary Size:** 7 product categories

Example:

```
prod2id = {
    'DISABILITY': 1,
    'HEALTH': 2,
    'INVESTMENT': 3,
    'LIFE_INSURANCE': 4,
    'NETWORK_PRODUCTS': 5,
    'OTHER': 6,
    'RETIREMENT': 7
}
```

3.2 Sequence Building and Deduplication

3.2.1 Group Events by Client

- Group all events by `cont_id`
- Sort events by `event_idx` within each client

- Extract product category sequence per client

3.2.2 Remove Consecutive Duplicates

Remove consecutive identical products in a sequence:

- **Rationale:** Focus on product transitions, not repeated purchases
- **Example:** [RETIREMENT, RETIREMENT, LIFE_INSURANCE] → [RETIREMENT, LIFE_INSURANCE]

3.2.3 Minimum Events Filter

Keep only clients with at least MIN_EVENTS (default: 2) products:

- Ensures at least 1 history item + 1 label per training example
- Result: ~388,326 users with sufficient history

3.3 Sliding Window Example Generation

For each client sequence, create multiple training examples using a sliding window approach:

Process:

1. For each position *i* in sequence (starting from index 1):
 - **History:** Last MAX_SEQ_LEN products before position *i*
 - **Label:** Product at position *i*
2. Convert products to numeric IDs using prod2id mapping
3. Pad history to MAX_SEQ_LEN if shorter (left-pad with 0)

Example:

- Sequence: [RETIREMENT, LIFE_INSURANCE, INVESTMENT] (IDs: [7, 4, 3])
- Examples generated:
 - History: [7], Label: 4 (LIFE_INSURANCE)
 - History: [7, 4], Label: 3 (INVESTMENT)

Result: ~580,620 training examples

3.4 History-Derived Features

For each training example, compute features from the history sequence:

3.4.1 Sequence Statistics

- seq_len: Length of history sequence (before padding)
- last_1: Product ID of most recent product in history
- last_2: Product ID of second most recent product (0 if sequence length < 2)
- unique_prior: Number of unique products in history
- num_switches: Number of product transitions in history

3.4.2 Frequency Features

For each product category (1 to NUM_CLASSES), compute:

- freq_{i}: Count of product ID *i* in history sequence

Example: If history contains [RETIREMENT, RETIREMENT, LIFE_INSURANCE] :

- freq_7 = 2 (RETIREMENT appears twice)
- freq_4 = 1 (LIFE_INSURANCE appears once)
- freq_1 through freq_6 = 0

Total Frequency Features: 7 (one per product category)

3.5 Padded History Sequence Features

Convert variable-length history sequences to fixed-length features:

- MAX_SEQ_LEN: 10 (maximum history length)
- Padding: Left-pad with 0s if sequence shorter than MAX_SEQ_LEN
- Truncation: Keep only last MAX_SEQ_LEN products if sequence longer

Features Created: hist_0, hist_1, ..., hist_9

- Each represents a position in the padded history sequence
- Value is product ID (0 for padding)

Example:

- History: [7, 4, 3] (length 3)
- Padded: [0, 0, 0, 0, 0, 0, 0, 7, 4, 3]
- Features: hist_0=0, hist_1=0, ..., hist_7=7, hist_8=4, hist_9=3

3.6 Static Client Features

Join with the most recent client snapshot (latest `register_ts` per client):

3.6.1 Financial Features

- `acct_val_amt`: Account value amount
- `face_amt`: Face amount
- `cash_val_amt`: Cash value amount
- `wc_total_assets`: Total assets
- `wc_assetmix_stocks`: Stock allocation
- `wc_assetmix_bonds`: Bond allocation
- `wc_assetmix_mutual_funds`: Mutual fund allocation
- `wc_assetmix_annuity`: Annuity allocation
- `wc_assetmix_deposits`: Deposit allocation
- `wc_assetmix_other_assets`: Other assets

3.6.2 Demographic Features

- `psn_age`: Person age

3.6.3 Categorical Features

- `client_seg`: Client segment
- `client_seg_1`: Client segment level 1
- `aum_band`: Assets under management band
- `channel`: Sales channel
- `agent_segment`: Agent segment
- `branchoffice_code`: Branch office code

Join Method: Left join to preserve all training examples

3.7 Missing Value Handling

3.7.1 Numeric Features

- Fill missing values with 0
- Applied to all numeric columns except: `label`, `seq_len`, `last_1`, `last_2`, `unique_prior`, `num_switches`

3.7.2 Categorical Features

- Compute mode (most frequent value) for each categorical column
- Fill missing values with mode
- If mode is null, use "UNKNOWN"

Categorical Columns: `client_seg`, `client_seg_1`, `aum_band`, `channel`, `agent_segment`, `branchoffice_code`

3.8 Categorical Encoding

Convert categorical string values to integer indices:

1. For each categorical column:
 - Extract all unique values
 - Sort alphabetically
 - Create mapping: value → index (0-based)
2. Create new column: `{categorical_col}_idx`
3. Map values using UDF (User Defined Function)

Example:

- channel values: ["Direct", "Advisor", "Online"]
- Mapping: {"Direct": 0, "Advisor": 1, "Online": 2}
- New column: `channel_idx` with integer values

Categorical Columns Encoded: 6 (each gets an `_idx` column)

3.9 Final Feature Set

Total Features: 39

Feature Categories:

1. **History Sequence Features** (10): `hist_0` through `hist_9`
2. **Sequence Statistics** (5): `seq_len`, `last_1`, `last_2`, `unique_prior`, `num_switches`
3. **Frequency Features** (7): `freq_1` through `freq_7`
4. **Financial Features** (11): Account values, face amount, cash value, total assets, asset mix (6 categories)
5. **Demographic Features** (1): `psn_age`

4. Sequence Building

4.1 Training Example Structure

Each training example represents a prediction task:

- **Input:** Client's purchase history up to a point in time
- **Output:** Next product category purchased

4.2 Example Generation Process

1. **Client Sequences:** Each client has a chronological sequence of product purchases
2. **Sliding Window:** For each position in the sequence (except the first):
 - Use all previous products as history
 - Use current product as label
3. **Multiple Examples per Client:** Clients with longer sequences generate more examples

Example Client Journey:

- Client purchases: [RETIREMENT, RETIREMENT, LIFE_INSURANCE, INVESTMENT]
- After deduplication: [RETIREMENT, LIFE_INSURANCE, INVESTMENT]
- Training examples:
 - Example 1: History= [RETIREMENT], Label= LIFE_INSURANCE
 - Example 2: History= [RETIREMENT, LIFE_INSURANCE], Label= INVESTMENT

4.3 Data Conversion

Convert from Spark DataFrame to Pandas DataFrame:

- **Reason:** LightGBM requires Pandas/NumPy input
- **Caching:** Cache Spark DataFrames before conversion for performance
- **Columns Selected:** cont_id, label, and all model_feature_cols

5. Train/Validation/Test Split

5.1 Split Strategy

- **Method:** Random split (stratified by label not used, but could be added)
- **Proportions:**
 - **Training:** 80% (TRAIN_FRAC = 0.8)
 - **Validation:** 10% (VAL_FRAC = 0.1)
 - **Test:** 10% (TEST_FRAC = 0.1)
- **Seed:** 42 (for reproducibility)

5.2 Split Implementation

```
train_spark, val_spark, test_spark = examples_full.randomSplit(  
    [TRAIN_FRAC, VAL_FRAC, TEST_FRAC],  
    seed=RANDOM_SEED  
)
```

5.3 Data Volumes

After split (approximate):

- **Training:** ~464,400 examples
- **Validation:** ~57,875 examples
- **Test:** ~58,345 examples

5.4 Label Remapping

Convert labels from product IDs (1-7) to zero-based indices (0-6) for LightGBM:

- Create `label_map`: Original product ID → Zero-based index
- Create new column `label0` in all splits
- Update `num_class` parameter to match actual number of classes

Example:

- Original labels: [1, 2, 3, 4, 5, 6, 7]
- Mapped labels: [0, 1, 2, 3, 4, 5, 6]
- num_classes : 7

5.5 Final Data Preparation

1. **Fill Missing Values:** Fill any remaining NaN values with 0
2. **Data Types:** Ensure all features are numeric (integers or floats)
3. **Feature Ordering:** Maintain consistent feature column order across all splits

6. Model Training

6.1 Model Algorithm

- **Algorithm:** LightGBM (Gradient Boosting Decision Trees)
- **Task:** Multiclass Classification
- **Number of Classes:** 7 product categories

6.2 Training Parameters

```
LGB_PARAMS = {  
    "objective": "multiclass",  
    "num_class": 7,  
    "metric": "multi_logloss",  
    "boosting_type": "gbdt",  
    "learning_rate": 0.05,  
    "num_leaves": 64,  
    "min_data_in_leaf": 50,  
    "feature_fraction": 0.8,  
    "subsample": 0.8,  
    "subsample_freq": 1,  
    "lambda_l2": 2.0,  
    "verbosity": -1  
}
```

Parameter Descriptions:

- **objective:** "multiclass" for multi-class classification
- **num_class:** Number of product categories (7)
- **metric:** "multi_logloss" (multi-class log loss)
- **learning_rate:** 0.05 (conservative learning rate)
- **num_leaves:** 64 (tree complexity)
- **min_data_in_leaf:** 50 (minimum samples per leaf)
- **feature_fraction:** 0.8 (use 80% of features per tree)
- **subsample:** 0.8 (use 80% of data per tree)
- **lambda_l2:** 2.0 (L2 regularization)

6.3 Training Configuration

- **Number of Boost Rounds:** 2000 (maximum iterations)
- **Early Stopping:** 50 rounds (stop if validation loss doesn't improve)
- **Validation Set:** Used for early stopping and monitoring

6.4 Training Process

1. **Create LightGBM Datasets:**
 - Training dataset: `train_pd[feature_cols_final]` with labels `train_pd["label0"]`
 - Validation dataset: `val_pd[feature_cols_final]` with labels `val_pd["label0"]`
2. **Train Model:**
 - Monitor both training and validation loss
 - Stop early if validation loss doesn't improve for 50 rounds
 - Save best model based on validation performance

8. Model Registration

8.1 MLflow Integration

The trained model is registered using MLflow for model versioning and deployment:

```
mlflow.lightgbm.log_model(  
    model,  
    artifact_path="final_lgbm_multiclass_model",  
    registered_model_name="eda_smartlist.models.final_lgbm_multiclass_model",  
    signature=signature,  
    input_example=input_example  
)
```

8.2 Model Artifacts

Registered Components:

1. **Model File:** LightGBM model binary
2. **Model Signature:** Input/output schema
3. **Input Example:** Sample input data for reference
4. **Metadata:** Training parameters, metrics, etc.

8.3 Model Version

Create an alias for easy model access:

- **Model Name:** eda_smartlist.models.final_lgbm_multiclass_model
- **Version:** 1

8.4 Artifact Storage

Additional artifacts saved separately:

- **artifacts.pkl:** Contains:
 - prod2id : Product to ID mapping
 - id2prod : ID to product mapping
 - label_map : Label remapping dictionary
 - num_classes : Number of classes
 - feature_cols : List of feature column names
 - max_seq_len : Maximum sequence length (10)
 - categorical_mappings : Categorical encoding mappings

Storage Location: /Workspace/Users/rajesh.patil@equitable.com/Final_model_files/

9. Key Parameters

9.1 Data Processing Parameters

Parameter	Value	Description
SAMPLE_FRACTION	0.2	Fraction of data to sample (None for full dataset)
MIN_EVENTS	2	Minimum events per client to create examples
MAX_SEQ_LEN	10	Maximum history length for features

9.2 Train/Test Split Parameters

Parameter	Value	Description
TRAIN_FRAC	0.8	Training set proportion (80%)
VAL_FRAC	0.1	Validation set proportion (10%)
TEST_FRAC	0.1	Test set proportion (10%)
RANDOM_SEED		Random seed for

Parameter	Value	Description
-----------	-------	-------------

9.3 Model Parameters

Parameter	Value	Description
NUM_BOOST_ROUND	2000	Maximum number of boosting rounds
EARLY_STOP	50	Early stopping rounds
learning_rate	0.05	Learning rate for gradient boosting
num_leaves	64	Maximum tree leaves
min_data_in_leaf	50	Minimum samples per leaf
feature_fraction	0.8	Feature sampling ratio
subsample	0.8	Data sampling ratio
lambda_12	2.0	L2 regularization

10. Data Flow Summary

10.1 Source to Training Examples

```

Source Table (48M+ rows)
↓
Product Category Transformation
↓
Filter: Active policies, non-null fields
↓
Optional Sampling (20%)
↓
Event Ordering (by register_date per client)
↓
Sequence Building (group by client)
↓
Deduplicate Consecutive Products
↓
Filter: Clients with ≥2 events
↓
Sliding Window Example Generation
↓
Feature Engineering
↓
Train/Val/Test Split (80/10/10)
↓
Training Examples (464K+ examples)

```

10.2 Feature Categories

1. **Temporal Features:** History sequence (hist_0 to hist_9)
2. **Sequence Statistics:** Length, last products, transitions
3. **Frequency Features:** Product purchase frequencies
4. **Financial Features:** Account values, asset allocations
5. **Demographic Features:** Age, client segments
6. **Categorical Features:** Channel, agent segment, branch office

11. Key Design Decisions

11.1 Why Sliding Window?

- **Captures Evolution:** Each example represents a point in the client's journey
- **Maximizes Data:** One client with N products generates N-1 examples
- **Temporal Context:** Model learns from sequences of different lengths

11.2 Why Remove Consecutive Duplicates?

- **Focus on Transitions:** Model predicts product changes, not repeated purchases
- **Reduces Noise:** Eliminates redundant information
- **Improves Signal:** Highlights meaningful product switches

11.3 Why LightGBM?

- **Handles Mixed Features:** Works well with numeric and categorical features
- **Efficiency:** Fast training on large datasets
- **Performance:** Strong results on tabular data
- **Interpretability:** Feature importance and SHAP values available

Appendix A: Column Reference

A.1 Source Table Columns Used

Identifiers:

- cont_id : Client/Contract ID
- axa_party_id : AXA Party ID
- branchoffice_code : Branch office code

Product Information:

- prod_lob : Product line of business
- sub_product_level_1 : Sub-product level 1
- sub_product_level_2 : Sub-product level 2
- register_date : Product registration date
- policy_status : Policy status (Active/Inactive)

Financial Metrics:

- acct_val_amt : Account value amount
- face_amt : Face amount
- cash_val_amt : Cash value amount
- wc_total_assets : Total assets
- wc_assetmix_stocks : Stock allocation
- wc_assetmix_bonds : Bond allocation
- wc_assetmix_mutual_funds : Mutual fund allocation
- wc_assetmix_annuity : Annuity allocation
- wc_assetmix_deposits : Deposit allocation
- wc_assetmix_other_assets : Other assets

Demographics:

- psn_age : Person age
- client_seg : Client segment
- client_seg_1 : Client segment level 1
- aum_band : Assets under management band

Channel Information:

- channel : Sales channel
- agent_segment : Agent segment
- division_name : Division name
- branch_name : Branch name
- business_city : Business city
- business_state_cod : Business state code

Temporal:

- business_month : Business month (YYYYMM format)

Appendix B: Product Category Mapping Rules

B.1 LIFE_INSURANCE

- prod_lob == "LIFE"
- sub_product_level_1 in: ["VLI", "WL", "UL/IUL", "TERM", "PROTECTIVE PRODUCT"]
- sub_product_level_2 contains "LIFE"
- sub_product_level_2 in: ["VARIABLE UNIVERSAL LIFE", "WHOLE LIFE", "UNIVERSAL LIFE", "INDEX UNIVERSAL LIFE", "TERM PRODUCT", "VARIABLE LIFE", "SURVIVORSHIP WHOLE LIFE", "MONY PROTECTIVE PRODUCT"]

B.2 RETIREMENT

- prod_lob in: ["GROUP RETIREMENT", "INDIVIDUAL RETIREMENT"]
- sub_product_level_1 in: ["EQUIVEST", "RETIREMENT 401K", "ACCUMULATOR", "RETIREMENT CORNERSTONE", "SCS", "INVESTMENT EDGE"]
- sub_product_level_2 contains: "403B", "401", "IRA", or "SEP"

B.3 INVESTMENT

- prod_lob == "BROKER DEALER"
- sub_product_level_1 in: ["INVESTMENT PRODUCT - DIRECT", "INVESTMENT PRODUCT - BROKERAGE", "INVESTMENT PRODUCT - ADVISORY", "DIRECT", "BROKERAGE", "ADVISORY", "CASH SOLICITOR"]
- sub_product_level_2 contains: "Investment", "Brokerage", or "Advisory"

B.4 NETWORK_PRODUCTS

- prod_lob == "NETWORK"
- sub_product_level_1 == "NETWORK PRODUCTS"
- sub_product_level_2 == "NETWORK PRODUCTS"

B.5 DISABILITY

- prod_lob == "OTHERS" AND sub_product_level_1 == "HAS"
- sub_product_level_2 == "HAS - DISABILITY"

B.6 HEALTH

- prod_lob == "OTHERS" (and not DISABILITY)
- sub_product_level_2 == "GROUP HEALTH PRODUCTS"

B.7 OTHER

- All products not matching above rules